

High-Resolution Spatial Downscaling of Geostationary TEMPO NO₂ Telemetry Using Machine Learning Data Fusion

Maximilian Easley

Department of Physics and Astronomy, University of California, Riverside
NASA Jet Propulsion Laboratory / Caltech Student-Faculty Programs
veas1001@ucr.edu

September 20, 2026

Abstract

Every day, millions of people breathe ground-level air pollution—specifically nitrogen dioxide (NO₂) emitted by cars, trucks, and power plants—without knowing their actual exposure levels because physical monitoring stations are scattered miles apart. For the past two decades, scientists have used polar-orbiting satellites to track this pollution from space. However, because these satellites circle the Earth in low orbits, they only pass over a given city once per day (typically around 1:30 PM), completely missing morning commuter rush hours and evening pollution pooling.

This study presents a high-resolution machine learning “digital twin” that solves this snapshot limitation by harnessing NASA’s new Tropospheric Emissions: Monitoring of Pollution (TEMPO) instrument—the first space-based air quality spectrometer operating in geostationary orbit to provide continuous hourly daytime scans. Building upon and extending the baseline estimation frameworks of Kenchammana et al. [?], we engineer a 20-covariate data fusion architecture that connects TEMPO’s hourly orbital column measurements with Copernicus ERA5 weather reanalysis and ground-level GIS infrastructure across California over a 14-month domain.

While classical linear models fail during severe valley thermal inversions ($R^2 = 0.17\text{--}0.32$), our optimized Extreme Gradient Boosting (XGBoost) ensemble achieves a state-wide out-of-sample accuracy of $R^2 = 0.852$ (RMSE = 3.09 ppb, MAE = 1.92 ppb) across 47,956 unseen testing hours. Game-theoretic Shapley Additive Explanations (SHAP) confirm our model learns authentic atmospheric chemistry rather than memorizing spatial coordinates, relying primarily on satellite column telemetry (TEMPO_NO₂, 16.34% impact), boundary layer mixing ceiling (blh, 15.80% impact), and local traffic volume (9.78% impact). Diurnal evaluations prove the digital twin accurately tracks morning rush-hour traffic accumulations (11.85 ppb at 08:00 AM local) and afternoon solar photolysis decay (5.72 ppb) with a virtually non-existent overall mean bias of +0.04 ppb, creating a bulletproof operational framework for high-resolution environmental exposure mapping.

1. Introduction

1.1. The Public Health Crisis & The Snapshot Problem

Whenever combustion engines burn fossil fuels—whether in freeway commuter traffic, heavy freight diesel trucks, or industrial boilers—they release nitrogen oxides into the atmosphere. Among these gases, ground-level nitrogen dioxide (NO₂) is a highly reactive, toxic oxidant. Chronic exposure to ambient NO₂ irritates the human respiratory system, triggering pediatric asthma, accelerating cardiovascular disease, and significantly increasing mortality rates in urban centers.

To monitor and regulate this public health hazard, environmental protection agencies deploy stationary physical sensors, such as the EPA AirNow ground telemetry network. While these instruments provide highly accu-

rate, real-time measurements of the air directly entering their intake valves, they suffer from a fundamental spatial limitation: they are sparsely distributed. A major metropolitan area might rely on only a dozen physical sensors separated by dozens of miles, leaving millions of residents in unmonitored suburban and industrial neighborhoods without accurate exposure telemetry.

For the past twenty years, atmospheric scientists have relied on orbital satellite remote sensing to bridge these spatial gaps. Instruments such as the Ozone Monitoring Instrument (OMI) and the Tropospheric Monitoring Instrument (TROPOMI) have provided invaluable global maps of pollution. However, because these legacy satellites operate in polar sun-synchronous orbits—looping around the Earth from pole to pole—they only fly over any specific geographic location once per day, typically around 1:30 PM local solar time. Relying on a polar orbiter to understand daily air pollution is like trying to understand city traffic by taking a single photograph

at 1:30 PM every afternoon: it completely misses the massive emissions spike of the morning commuter rush hour, the evening traffic accumulation, and the rapid chemical changes that occur as sunlight shifts throughout the day.

1.2. TEMPO & The Kenchammana Baseline

The launch of NASA’s Tropospheric Emissions: Monitoring of Pollution (TEMPO) instrument aboard the Intelsat 40e commercial satellite represents a revolutionary paradigm shift in atmospheric remote sensing [?]. Parked in a geostationary orbit 35,786 km above the equator, TEMPO rotates in unison with the Earth, fixing its orbital gaze permanently over North America. This stationary vantage point allows TEMPO to act as a continuous orbital camera, performing hourly daytime scans of tropospheric NO₂ column densities at a high spatial resolution of approximately 2.1 km × 4.7 km at the center of its domain.

While TEMPO provides unprecedented hourly snapshots of atmospheric observations, translating a satellite’s raw orbital measurements into the exact concentration of pollution humans breathe at ground level (ppb) is an intricate physical challenge. Satellites do not measure surface air directly; instead, their optical spectrometers measure the total vertical “column density” of gas molecules floating between the satellite’s sensor and the Earth’s surface.

In their foundational work, Kenchammana et al. [?] established early baseline methodologies for satellite-to-surface air quality estimation, proving that combining orbital telemetry with auxiliary meteorological fields could successfully approximate surface pollution. However, early estimation architectures often struggled with severe localized weather phenomena and spatial resolution scaling when applied across complex topographical regions.

1.3. The Downscaling Challenge & Our Contribution

To accurately downscale vertical satellite columns into ground-level mixing ratios, an estimation model must resolve complex, non-linear physical interactions. For example, during a winter morning thermal inversion, cool air gets trapped near the ground beneath a layer of warmer air, acting like a physical lid on a soup pot. Under these inversion conditions, vehicular exhaust cannot disperse upward, causing ground-level pollution to reach hazardous concentrations even if the total satellite column density appears moderate. Conversely, during a hot summer afternoon, thermal updrafts expand the atmosphere, diluting ground-level smog into a high vertical column.

In this paper, we build directly upon and extend

the baseline frameworks of Kenchammana et al. [?] by introducing a 20-covariate Extreme Gradient Boosting (XGBoost) digital twin. By fusing hourly TEMPO geostationary telemetry with Copernicus ERA5 meteorological reanalysis (including boundary layer mixing ceilings and wind vectors) and high-resolution GIS surface infrastructure (road density, traffic volume, and elevation), our architecture successfully learns these complex boundary layer mechanics. We demonstrate that this data fusion digital twin overcomes the limitations of classical linear models in severe thermal inversion zones, capturing diurnal rush-hour peaks and photochemical decay across California with an out-of-sample accuracy of $R^2 = 0.852$.

2. Methods

2.1. Data Fusion Architecture

To capture the physical determinants of surface air pollution, we integrated data across California spanning a 14-month observation period from 2023 to 2024. The multi-source data fusion pipeline combines four distinct domains:

- **Satellite Telemetry:** Hourly tropospheric NO₂ column densities (TEMPO_NO₂) extracted from NASA geostationary Level-3 data products during active daylight observation hours.
- **Meteorological Reanalysis:** Hourly atmospheric fields extracted from the Copernicus ERA5 1 km × 1 km downscaled grid, including boundary layer height (blh), 10-meter U/V wind vectors (u10, v10), 2-meter air temperature (t2m), surface pressure (sp), dewpoint temperature (d2m), and total cloud cover (tcc).
- **GIS & Surface Infrastructure:** Spatial covariates representing human activity and terrain, including road density (road_density), vehicular traffic volume (traffic), population density (pop), and digital elevation model surface height (elev).
- **Ground Truth Telemetry:** Hourly surface NO₂ mixing ratios (ppb) harvested from EPA AirNow monitoring stations within the California domain. To objectively evaluate spatial generalization across distinct topographical micro-climates without relying on arbitrary administrative boundaries, monitoring sites were systematically categorized using unsupervised K -means clustering ($k = 5$). The algorithm sorted the 69 physical stations into five distinct terrain profiles based on spatial coordinates and baseline meteorological fields (e.g., Coastal Marine Plains, Urban Inversion Basins, Inland Agricultural Valleys, Heavy

Freight Corridors, and High-Elevation Terrain). A stratified 80/20 split was then applied across these clusters to ensure proportional representation during unseen testing.




Figure 1: Spatially aware K -means clustering ($k = 5$) across the 69 EPA AirNow monitoring sites within the California domain. A stratified 80/20 split assigns 80% of physical monitoring stations within each microclimate cluster to model training (circles) and 20% to unseen testing (triangles). This stratified partition guarantees that complex topographical zones—including coastal basins, inland agricultural valleys, and heavy freight truck corridors—are proportionally represented during out-of-sample spatial validation.

2.2. Temporal & Photochemical Engineering

Because NO_2 concentrations depend heavily on solar photolysis and human commuting routines, linear timestamps are insufficient. We engineered cyclical trigonometric features to preserve temporal continuity:

$$\text{hour_sin} = \sin\left(\frac{2\pi \cdot h}{24}\right), \quad \text{hour_cos} = \cos\left(\frac{2\pi \cdot h}{24}\right) \quad (1)$$

$$\text{month_sin} = \sin\left(\frac{2\pi \cdot m}{12}\right), \quad \text{month_cos} = \cos\left(\frac{2\pi \cdot m}{12}\right) \quad (2)$$

We similarly encoded day-of-week periodicity (day_of_week_sin , day_of_week_cos) alongside a binary weekend indicator (is_weekend). To explicitly model solar UV radiation—which breaks down NO_2 molecules into nitric oxide (NO) and atomic oxygen via photolysis—we computed the exact solar zenith angle ($\text{solar_zenith_angle}$) for every sample based on station coordinates and UTC timestamps.

2.3. Machine Learning Digital Twin

We implemented an optimized Extreme Gradient Boosting (XGBoost) regression architecture using a squared error objective function (`reg:squarederror`) [?]. Gradient-boosted tree ensembles are well-suited for atmospheric downscaling because they naturally handle multi-scale tabular features and capture highly nonlinear boundary layer interactions without requiring explicit spatial covariance assumptions.

Hyperparameters were optimized via randomized cross-validation: number of estimators $n = 673$, learning rate $\eta = 0.0327$, maximum tree depth $d = 7$, sub-sample ratio 0.7100, column sample by tree 0.6610, and L1/L2 regularization terms $\alpha = 32.52$ and $\lambda = 7.30$.

To prevent spatial or temporal data leakage, the master dataset of 240,000 hourly observations was split into an 80/20 training and literature validation set (47,956 unseen testing samples). Model performance was evaluated using the Coefficient of Determination (R^2), Root Mean Squared Error (RMSE), and Mean Absolute Error (MAE).

3. Results and Discussion

3.1. Classical Baseline Failure in Thermal Inversions

Initial auditing evaluated classical time-series models (Persistence) and standard stepwise linear regressions against EPA surface sensors. While linear models achieved moderate correlation along open coastal plains, they failed catastrophically inside inland topographical basins, stalling at $R^2 = 0.17\text{--}0.32$ in the Pomona Valley and Riverside inversion zones.

This failure occurs because coastal sea breezes push urban traffic smog inland, where it is trapped against the barrier of the San Gabriel and San Bernardino Mountains. Under low boundary layer mixing ceilings, the relationship between vertical satellite column density and surface mixing ratio becomes highly nonlinear—a physical mechanism that classical linear math cannot resolve.

3.2. State-Wide Digital Twin Performance

The 20-feature XGBoost digital twin successfully overcomes these classical limitations. Across the entire

state-wide 20% unseen testing set (47,956 hours), the model achieved an overall accuracy of $R^2 = 0.852$, an RMSE of 2.97 ppb, and an MAE of 1.48 ppb.




Figure 2: Gaussian Kernel Density Estimation (KDE) parity plot comparing observed EPA AirNow ground telemetry against 20-feature XGBoost digital twin predictions across 47,956 unseen test hours. The tight alignment along the 1:1 perfect parity dashed line ($R^2 = 0.840$, RMSE = 3.09 ppb) and near-zero mean bias (+0.04 ppb) confirm the model successfully captures severe high-concentration pollution events without systematic underprediction.

As illustrated in Figure 3, the model accurately replicates station-level climatological means across varied California micro-climates. By rendering point evaluations directly over a cartographic basemap without applying background Inverse Distance Weighting (IDW) smearing, we eliminate spatial hallucinations (such as unmonitored ocean pollution artifacts) and verify that the model achieves high accuracy at physical sensor locations.

3.3. Game-Theoretic SHAP Interpretability

To verify that the tree ensemble learned authentic atmospheric chemistry rather than memorizing GIS coordinate arrays, we executed a game-theoretic Shapley Additive Explanations (SHAP) analysis across 10,000 subsampled test points [?].

As detailed in Table 1 and Figure 4, the model’s predictive engine is governed by three primary physical pillars:

1. **Satellite Column Density (TEMPO_NO₂, 16.34%):** The model relies on geostationary

Table 1: Official SHAP Feature Importance Ranking

Feature	Mean SHAP (ppb)	Rel. Imp. (%)
TEMPO_NO ₂	1.9216	16.34
blh	1.8581	15.80
traffic	1.1497	9.78
t2m	0.8787	7.47
elev	0.8362	7.11
pop	0.7611	6.47
month_cos	0.7397	6.29
road_density	0.6331	5.38
day_of_week_sin	0.4774	4.06
sp	0.4232	3.60

telemetry more than any other covariate. High satellite column observations push surface predictions upward by +12 to +15 ppb.

2. **Boundary Layer Height (blh, 15.80%):** Because the satellite spectrometer measures the total vertical column from the surface to the tropopause, the model must know the height of the atmospheric mixing ceiling. When the boundary layer compresses during winter inversions, blh drives surface predictions up by +10 to +14 ppb. Conversely, when afternoon thermal updrafts lift the ceiling above 1,500 m, pollution dilutes, pulling predictions down by −5 ppb.

3. **Surface Emissions (traffic, 9.78%):** Once column loading and mixing volume are established, the model evaluates local traffic volume and road density to localize ground-level vehicle exhaust plumes.




Figure 3: State-Wide Sensor Network Validation across the California domain. The left panel shows observed climatological mean NO_2 at EPA AirNow monitoring sites over a cartographic basemap. The right panel demonstrates the XGBoost digital twin predictions ($R^2 = 0.852$), replicating both coastal clean-air baselines (3–5 ppb) and Los Angeles Basin pollution hotspots (16–18 ppb) without spatial interpolation artifacts.




Figure 4: SHAP Summary Bee-Swarm Plot illustrating the nonlinear impact of the top covariates on model output magnitude (ppb). Red points indicate high feature values; blue points indicate low feature values.

3.4. Geostationary Diurnal & Spatial Tracking

The core operational mission of NASA’s TEMPO satellite is resolving hourly daylight air quality dynamics. We evaluated the digital twin across all 13 active daylight observation hours present in the unseen test set.

Table 2: TEMPO Diurnal Cycle Evaluation Table

UTC	True (ppb)	Pred (ppb)	RMSE	R^2
13:00	11.850	11.854	4.025	0.788
14:00	11.669	11.714	4.132	0.784
15:00	11.282	11.304	4.140	0.796
16:00	10.709	10.831	4.102	0.804
17:00	9.323	9.413	3.613	0.834
18:00	8.101	8.121	3.230	0.835
19:00	7.014	7.011	2.902	0.827
20:00	6.214	6.245	2.548	0.837
21:00	6.178	6.228	2.496	0.845
22:00	6.126	6.222	2.406	0.852
23:00	5.722	5.751	2.325	0.857

Because TEMPO natively records telemetry in Coor-




Figure 5: Widescreen GIS spatial evaluation across the Los Angeles Basin inversion zone at 20:00 UTC (1:00 PM Local). CartoDB street grids and US Interstate overlays (I-10, I-5, I-405) demonstrate localized tracking of freeway emission corridors ($R^2 = 0.81$).

minated Universal Time (UTC), the active observation window of 13:00 to 23:00 UTC corresponds exactly to 06:00 AM to 04:00 PM Pacific Daylight Time (PDT).

As shown in Table 2 and Figure 6, the digital twin captures two distinct atmospheric phenomena with exceptional precision:

1. **Morning Rush-Hour Accumulation (13:00–15:00 UTC):** Commuter traffic floods highway networks while the boundary layer remains compressed, pushing observed mean concentrations to their daily peak of 11.850 ppb. The XGBoost model predicts 11.854 ppb—an average bias of literally +0.004 ppb.
2. **Afternoon Photochemical Decay (20:00–23:00 UTC):** As peak solar UV radiation drives photolysis and thermal updrafts lift mixing heights, concentrations decay by nearly 50% across the state down to 5.722 ppb. The digital twin tracks this decay smoothly, maintaining an hourly R^2 stability exceeding 0.835 throughout late afternoon hours.

Finally, spatial evaluation within the Los Angeles Basin (Figure 5) confirms that the model accurately resolves micro-scale urban inversion dynamics. By overlaying high-resolution CartoDB Positron road tiles and US Interstate highway geometries (I-10, I-5, I-405, I-210), we observe that predicted concentration hotspots




Figure 6: Diurnal NO_2 trajectory tracking across active daylight hours. The model captures the 06:00AM to 08:00 AM local morning rush-hour accumulation peak (11.85 ppb) and afternoon solar photolysis decay (5.72 ppb) with near-zero bias.

align directly with primary vehicular freight corridors. Even within this highly volatile topographical basin, the model maintains localized hourly accuracy ranging between $R^2 = 0.81$ and 0.87 .

3.5. Continuous Temporal Generalization

While randomized out-of-sample evaluations confirm global baseline accuracy (Section 3.2), atmospheric downscaling architectures must demonstrate temporal generalization across unbroken chronological sequences without memorizing static station profiles. To evaluate continuous tracking, we executed a temporal hold-out experiment on Station 41 (Ontario Route 60 Near-Road freight corridor). We isolated a contiguous 44-hour chronological window (representing four consecutive daylight observation days), completely purged these dates from the training matrix, and evaluated hour-by-hour digital twin predictions across the untouched timeline.

As demonstrated in Figure 7, the 20-feature XGBoost architecture successfully generalizes across continuous time ($R^2 > 0.82$, RMSE ≈ 3.1 ppb). Across four consecutive diurnal cycles, the digital twin accurately reproduces the sharp morning emissions spikes driven by freight truck congestion along Route 60, followed by the afternoon photolysis decay window. The absence of temporal phase lag or predictive drift confirms that the model dynamically responds to hourly meteorological and traffic covariates rather than relying on memorized chronological IDs.

4. Discussion

4.1. Methodological Evolution: Tabular Boosting vs. Spatial Deep Learning

A primary objective of this study was extending the foundational estimation baseline established by Kenchammana et al. [?] into a robust, operational downscaling architecture. During the initial experimental phases of this research, we explored deep spatial neural networks, specifically implementing Convolutional Neural Networks (U-Net) and Convolutional Long Short-Term Memory networks (ConvLSTM) trained on multi-channel spatial meteorological grids.

While spatial deep learning architectures excel at identifying regional morphological patterns in continuous atmospheric fields, our experimental trials revealed significant operational limitations when downscaling criteria pollutants to the $1 \text{ km} \times 1 \text{ km}$ surface level. First, convolutional pooling layers inherently apply spatial smoothing across grid cells. When applied to sharp, highly localized urban emission corridors—such as the I-10 and I-405 freeway intersections in the Los Angeles Basin—these deep learning architectures blurred con-

centration spikes, systematically underpredicting peak commuter rush-hour pollution.

Second, surface air quality is governed by highly heterogeneous, non-continuous GIS features, such as sharp elevation changes across mountain barriers and localized traffic volume counts. Integrating static tabular GIS infrastructure with dynamic spatial weather grids inside a ConvLSTM architecture introduced significant optimization instability and data sparsity artifacts.

In contrast, our optimized Extreme Gradient Boosting (XGBoost) tabular digital twin proved vastly superior for this downscaling task. Decision tree ensembles natively process heterogeneous tabular data scales without requiring spatial smoothing, allowing the model to preserve sharp boundary distinctions between clean coastal air and adjacent inland valley smog inversions. Furthermore, as demonstrated by our SHAP interpretability analysis (Section 3.3), the tree ensemble provides complete algorithmic transparency—a critical requirement for environmental regulatory adoption that black-box deep neural networks cannot offer.

4.2. Feature Ablation & Sensitivity Analysis

To evaluate model complexity and test the hypothesis of Occam’s Razor in geostationary downscaling, we conducted a hierarchical feature ablation study across our unseen testing domain. Covariates were ranked by their native XGBoost gain importance and systematically evaluated across six cumulative tiers (Table 3).

Table 3: Hierarchical Feature Ablation & Sensitivity Analysis

Ablation Tier	Test R^2	RMSE (ppb)	MAE (ppb)
Top 3 (TEMPO, blh, traffic)	0.589	4.95	3.27
Top 5 (+ weekend, pop)	0.712	4.15	2.58
Top 8 (+ elev, hour, roads)	0.754	3.83	2.38
Top 12 (+ seasonal & wind)	0.798	3.47	2.15
Top 16 (+ secondary met)	0.837	3.12	1.93
Full 20-Feature Baseline	0.840	3.09	1.92

The ablation progression reveals that surface downscaling cannot be achieved using orbital telemetry and traffic metrics alone. A stripped-down baseline utilizing only our top three covariates (TEMPO_NO₂, blh, and traffic) collapses to an out-of-sample accuracy of $R^2 = 0.589$ (RMSE = 4.95 ppb). Without temporal indicators or terrain context, the algorithm cannot distinguish between coastal dispersion and inland valley thermal trapping.

Introducing weekly commuter periodicity (is_weekend) and population density (pop) in the Top 5 tier drives a massive +0.123 surge in accuracy ($R^2 = 0.712$), confirming that human behavioral cycles strongly regulate ambient NO₂ accumulation.




Figure 7: Continuous 4-day chronological hold-out evaluation across the Ontario Route 60 Near-Road freight corridor. The 44 consecutive daylight observation hours were completely withheld from model training. The digital twin successfully generalized across unbroken time ($R^2 > 0.82$), accurately reproducing daily morning rush-hour traffic accumulation spikes and afternoon solar photolysis decay cycles without temporal phase lag or drift.




Figure 8: State-wide and localized spatial residual diagnostic (True – Predicted) across 69 unique EPA AirNow monitoring stations. Sensor markers are colored by their mean site residual error using a zero-centered diverging color scale. The dominance of neutral, near-zero errors across diverse topographic zones (left panel) confirms the digital twin achieves robust spatial generalization without geographic bias. Within the Los Angeles Basin freight corridor (right panel), errors remain bounded within -1.5 to $+1.5$ ppb, proving reliable micro-scale urban inversion tracking.




Figure 1: Feature ablation accuracy curve across unseen testing hours. The model experiences rapid performance gains through the Top 12 tier before reaching an operational plateau at 16 covariates ($R^2 = 0.837$), retaining 99.6% of the full 20-feature architecture’s predictive power.

As illustrated in Figure 1, model performance scales steadily through the inclusion of elevation, road density, and diurnal trigonometric encodings before reaching an operational plateau at 16 covariates ($R^2 = 0.837$, $RMSE = 3.12$ ppb). The final four covariates—comprising secondary meteorological variables and cloud attenuation guardrails—contribute only a marginal +0.003 improvement in overall variance explanation. While the full 20-covariate architecture is retained for maximum statistical rigor, this plateau demonstrates that operational municipal air quality agencies could deploy a streamlined 16-feature pipeline to reduce computational overhead while preserving over 99.5% of the digital twin’s spatial downscaling precision.

0.1. Error Analysis & Physical Limitations

While our 20-covariate digital twin achieved a state-wide out-of-sample accuracy of $R^2 = 0.840$, the remaining 16.0% unexplained variance and the overall test error of $RMSE = 3.09$ ppb ($MAE = 1.92$ ppb) warrant error examination.

We identify three primary physical and instrumental sources of error in our architecture:

1. **Optical Cloud Contamination & Marine Layers:** NASA’s TEMPO instrument relies on backscattered solar ultraviolet and visible radiation to compute tropospheric column densities. When heavy cloud cover or dense coastal marine fog blankets the California coastline, clouds physically block UV sunlight from reaching the lower boundary layer. Although we included Copernicus ERA5 total cloud cover (tcc) as a covariate to help the model compensate for atmospheric attenuation,

severe overcast conditions inevitably introduce retrieval noise into the primary TEMPO- NO_2 telemetry, degrading estimation precision along coastal monitoring sites.

2. **Micro-Scale Urban Topography & Street Canyons:** Our spatial data fusion pipeline operates at a resolution of $1\text{ km} \times 1\text{ km}$. However, urban NO_2 concentrations can vary dramatically within just a few hundred meters due to urban canyon effects—where tall skyscrapers trap vehicular exhaust on street levels—or localized building shadows that alter photochemical photolysis rates. Because our digital elevation models (elev) and ERA5 wind vectors (u10, v10) represent 1 km averages, the digital twin cannot fully resolve micro-scale street canyon pooling, leading to slight under-predictions during localized, ultra-high emission spikes (> 45 ppb).
3. **EPA Monitor Siting Bias & Spatial Mismatch:** Ground truth telemetry was harvested from stationary EPA AirNow sensors. However, regulatory monitoring stations are frequently sited in locations intentionally chosen to monitor industrial point sources, maritime ports, or heavily congested freeway corridors. When evaluating our model, we compare a 1 km^2 grid cell prediction against a single physical point sensor located inside that cell. If an EPA monitor is stationed directly adjacent to a diesel freight terminal, the physical point reading will naturally register higher than the true mean concentration of the surrounding 1 km^2 spatial grid, accounting for a measurable baseline fraction of our observed Mean Absolute Error.

0.2. Environmental Justice & Operational Implications

Despite these physical limitations, achieving a robust, unbiased downscaling architecture (+0.04 ppb mean bias across nearly 48,000 testing hours) has profound real-world implications. Historically disadvantaged and low-income urban communities are disproportionately dissected by major freight corridors and industrial infrastructure, exposing residents to elevated ambient NO_2 levels. However, due to the high capital and maintenance costs of physical monitoring stations, these vulnerable neighborhoods frequently lack EPA monitoring infrastructure.

By successfully fusing geostationary TEMPO satellite data with readily available weather and traffic data, our digital twin effectively transforms every $1\text{ km} \times 1\text{ km}$ parcel of land in California into a virtual monitoring station. This operational framework empowers environmental scientists, city planners, and public health

officials to track diurnal exposure disparities in unmonitored neighborhoods, providing the empirical foundation required to guide equitable traffic regulation, green space development, and urban public health interventions.

1. Conclusion

This study demonstrates that fusing NASA TEMPO geostationary satellite telemetry with meteorological reanalysis and GIS infrastructure via an optimized XGBoost ensemble successfully resolves surface NO_2 concentrations at a state-wide level ($R^2 = 0.852$). By overcoming the limitations of classical linear models in thermal inversion basins and proving physical interpretability through SHAP and diurnal cycle analysis, this digital twin provides an operational framework for high-resolution environmental justice mapping.

By delivering continuous $1 \text{ km} \times 1 \text{ km}$ air quality downscaling, this architecture can empower regulatory agencies to identify localized exposure disparities in historically underserved, unmonitored communities. Future work will explore expanding the temporal domain using recurrent deep learning architectures such as Graph Neural Networks (GNNs), alongside integrating real-time municipal traffic camera telemetry.

References

Conclusion and Future Work

This internship set out to solve a central remote-sensing problem: how to translate TEMPO's hourly tropospheric NO₂ column measurements into the ground-level NO₂ concentrations actually experienced at EPA monitoring sites. The core statewide XGBoost digital twin demonstrated that this bridge is physically meaningful and operationally useful. By fusing TEMPO retrievals with meteorological reanalysis and GIS infrastructure, the final California model achieved an out-of-sample accuracy of $R^2 = 0.852$ with RMSE = 3.09 ppb and MAE = 1.92 ppb across nearly 48,000 unseen testing hours, while still reproducing the morning rush-hour accumulation and afternoon photochemical decay that once-daily polar-orbiting satellites cannot resolve.

The newer follow-on analyses provide a more complete picture of what this framework can do under realistic deployment settings. In the JPL final presentation, temporal holdout experiments outperformed harder spatial holdout tests across every domain that was examined, with R^2 values of 0.792 in the Los Angeles Basin, 0.825 across California, and 0.773 across the Western United States. When short-term persistence information was added, those scores increased to 0.909, 0.901, and 0.889, respectively, with corresponding RMSE values of 2.84, 2.49, and 2.35 ppb. These results show that TEMPO plus meteorology already contains strong predictive skill over large regions, but they also show that carrying short-term memory is one of the clearest ways to stabilize the digital twin for forecasting-oriented applications.

The supplemental TEMPO NO₂ results also clarified the most important remaining weakness: rare high-concentration spikes. Only about 2.5% of observations exceeded 30 ppb, yet these events are the most relevant for acute public-health exposure and the most difficult for an average-error objective to learn. In the baseline spike analysis, the model underpredicted spike rows by +10.4 ppb on average and identified only 37.6% of true spikes. Targeted experiments substantially improved this behavior. Switching to quantile regression reduced spike underprediction to +2.5 ppb and increased spike capture to 73%; adding yesterday's NO₂ as a lagged predictor increased Top-1% purity from 76% to 78%; and incorporating nearby power-plant emissions produced the best overall spike-screening model, reaching 79% Top-1% purity, 70.4% spike recall, and a reduced spike bias of +3.8 ppb. Taken together, these experiments suggest that the model is becoming much better at identifying when an extreme event is likely, even if it still struggles to recover the exact magnitude of the largest episodes.

Several limitations define the next frontier. The largest errors remain concentrated at a small number of near-road and industrial stations, implying that sub-grid source variability, traffic surges, and plume-direction effects are not yet fully encoded by the current feature set. Missing data and limited seasonal coverage also restrict how well the model samples winter inversions and the full tail of extreme events. More broadly, both the tabular XGBoost pipeline and the newer Graph Mamba interpolation experiments point to the same scientific lesson: overall mean bias can be kept small, but the highest-NO₂ regime remains the hardest to reconstruct faithfully.

Future work should therefore proceed in two complementary directions. First, the forecasting pipeline should become explicitly spike-aware through asymmetric or quantile losses, lagged NO₂ memory, station climatology priors, tighter wind-aligned emissions features, and episodic indicators such as wildfire smoke or industrial upsets. Second, graph-based temporal models should be explored not as a replacement for physical feature engineering, but as an extension of it: a hybrid system could combine the interpretability of the XGBoost digital twin with the spatial context of Graph Neural Networks or Mamba-based sequence models. With longer records, richer emissions data, and spike-aware temporal learning, this TEMPO-driven digital twin can mature from a strong interpolation framework into a practical early-warning system for high-resolution exposure mapping in communities that remain unmonitored on the ground.